

# Rain Prediction

Import Libraries

```
In [2]: import matplotlib.pyplot as plt
import seaborn as sns
import datetime
from sklearn.preprocessing import LabelEncoder
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import seaborn as sns
from keras.layers import Dense, BatchNormalization, Dropout, LSTM
from keras.models import Sequential
from keras.utils import to_categorical
from keras.optimizers import Adam
from tensorflow.keras import regularizers
from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report, accuracy_score, f1_
from keras import callbacks

np.random.seed(23)
```

Load Data

```
In [3]: data = pd.read_csv("../input/weather-dataset-rattle-package/weatherAUS.csv")
data.head()
```

```
Out[3]:   Date  Location  MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustDir  WindGustSpeed  WindDir9am  ...  Humidity9am  Hu
0  2008-12-01    Albury     13.4      22.9       0.6        NaN        NaN          W         44.0           W  ...          71.0
1  2008-12-02    Albury      7.4      25.1       0.0        NaN        NaN         WNW         44.0          NNW  ...          44.0
2  2008-12-03    Albury     12.9      25.7       0.0        NaN        NaN         WSW         46.0           W  ...          38.0
3  2008-12-04    Albury      9.2      28.0       0.0        NaN        NaN          NE         24.0           SE  ...          45.0
4  2008-12-05    Albury     17.5      32.3       1.0        NaN        NaN          W         41.0          ENE  ...          82.0
```

5 rows × 23 columns

## About the data:

The dataset contains approximately 10 years of daily weather observations from various locations across Australia. The observations were drawn from numerous weather stations.

In this project, I will use this data to predict whether or not it will rain the next day. There are 23 attributes, including the target variable "RainTomorrow," indicating whether it will rain the following day or not.

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Date              145460 non-null   object 
 1   Location          145460 non-null   object 
 2   MinTemp           143975 non-null   float64
 3   MaxTemp           144199 non-null   float64
 4   Rainfall          142199 non-null   float64
 5   Evaporation       82670 non-null   float64
 6   Sunshine          75625 non-null   float64
 7   WindGustDir       135134 non-null   object 
 8   WindGustSpeed     135197 non-null   float64
 9   WindDir9am        134894 non-null   object 
 10  WindDir3pm        141232 non-null   object 
 11  WindSpeed9am     143693 non-null   float64
 12  WindSpeed3pm     142398 non-null   float64
 13  Humidity9am       142806 non-null   float64
 14  Humidity3pm       140953 non-null   float64
 15  Pressure9am      130395 non-null   float64
 16  Pressure3pm       130432 non-null   float64
 17  Cloud9am          89572 non-null   float64
 18  Cloud3pm          86102 non-null   float64
 19  Temp9am           143693 non-null   float64
 20  Temp3pm           141851 non-null   float64
 21  RainToday          142199 non-null   object 
 22  RainTomorrow       142193 non-null   object 

dtypes: float64(16), object(7)
memory usage: 25.5+ MB
```

The dataset contains missing values.

The dataset comprises both numeric and categorical values

#### Data Cleaning and visualization

Generate a count plot of the target column.

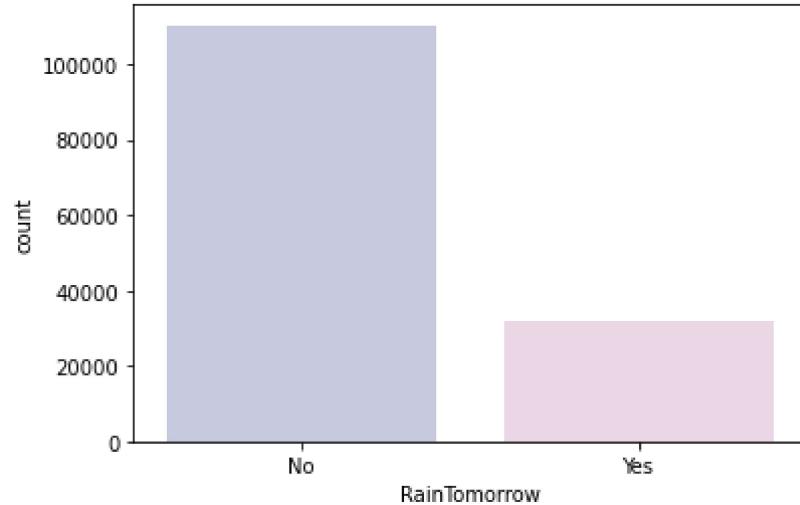
Calculate the correlation among numeric attributes.

Parse dates into the datetime format.

Encode days and months as continuous cyclic features.

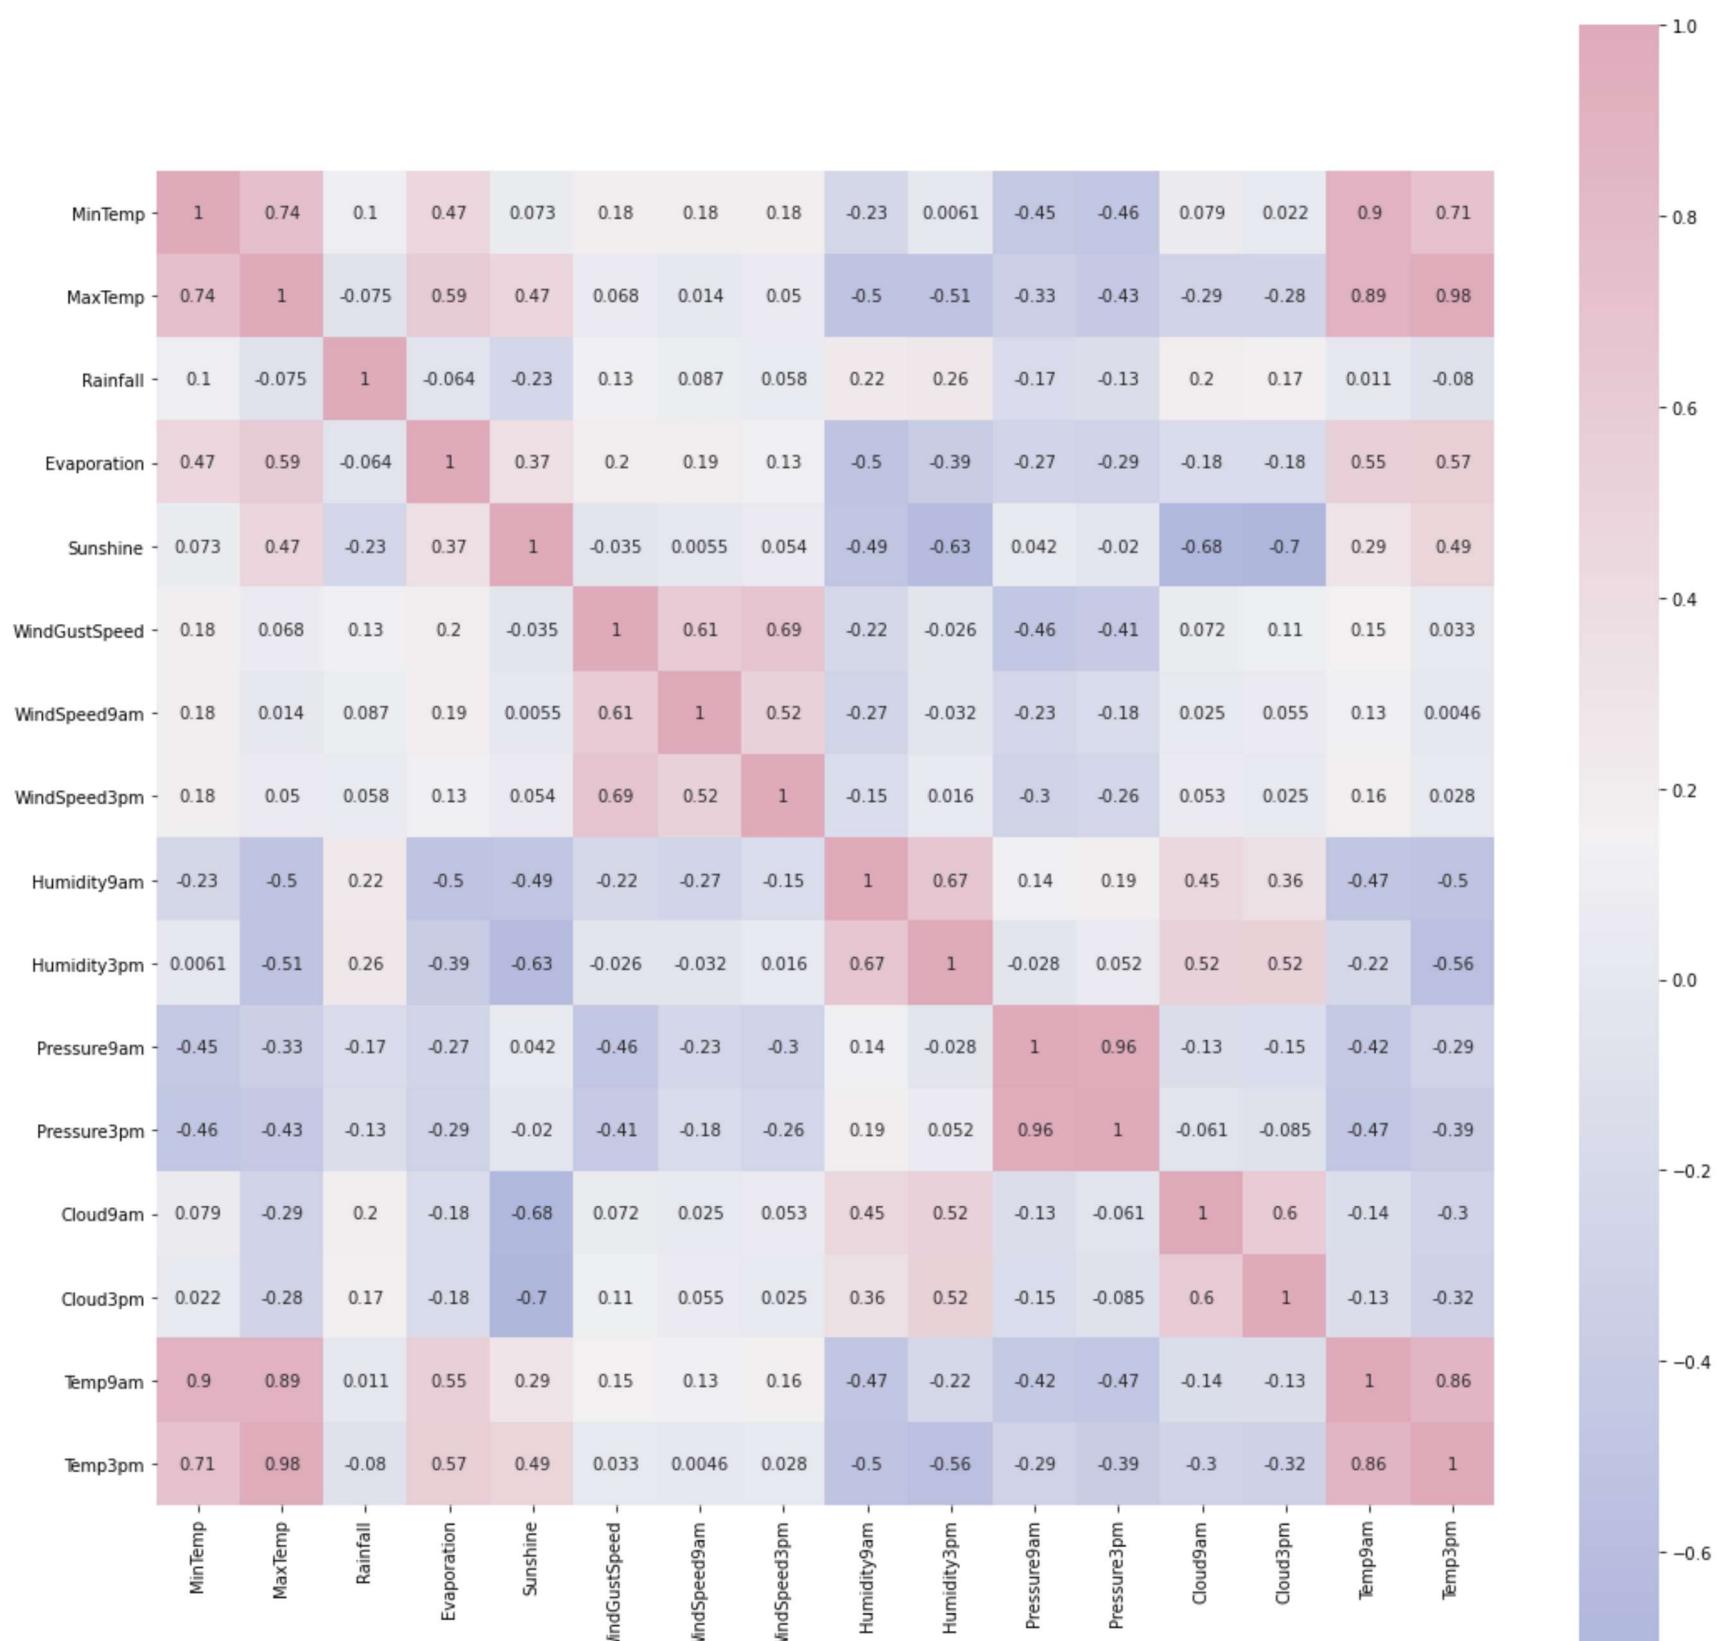
```
In [5]: # Evaluate the target and find out if data is imbalanced or not
cols= ["#C2C4E2", "#EED4E5"]
sns.countplot(x= data["RainTomorrow"], palette= cols)
```

```
Out[5]: <AxesSubplot:xlabel='RainTomorrow', ylabel='count'>
```



```
In [6]: # Correlation amongst numeric attributes
corrmat = data.corr()
cmap = sns.diverging_palette(260, -10, s=50, l=75, n=6, as_cmap=True)
plt.subplots(figsize=(18,18))
sns.heatmap(corrmat, cmap= cmap, annot=True, square=True)
```

```
Out[6]: <AxesSubplot:>
```



## Parse Dates into datetime.

My goal is to construct an artificial neural network (ANN). I will appropriately encode dates by representing months and days as cyclic continuous features. Since date and time exhibit inherent cyclic patterns, I divide them into periodic subsections, specifically years, months, and days. For each of these subsections, I generate two new features by calculating the sine and cosine transforms of the corresponding time units.

```
In [7]: #Parsing datetime
#exploring the Length of date objects
lengths = data["Date"].str.len()
lengths.value_counts()
```

```
Out[7]: 10    145460
Name: Date, dtype: int64
```

```
In [8]: #There don't seem to be any error in dates, so parsing values into datetime
data['Date']= pd.to_datetime(data["Date"])
#Creating a column of year
data['year'] = data.Date.dt.year

# Function for encoding datetime into cyclic parameters.
#As I intend to utilize this data in a neural network, I favor representing months and days as cyclic continuous features.
def encode(data, col, max_val):
    data[col + '_sin'] = np.sin(2 * np.pi * data[col]/max_val)
    data[col + '_cos'] = np.cos(2 * np.pi * data[col]/max_val)
    return data

data['month'] = data.Date.dt.month
data = encode(data, 'month', 12)

data['day'] = data.Date.dt.day
data = encode(data, 'day', 31)

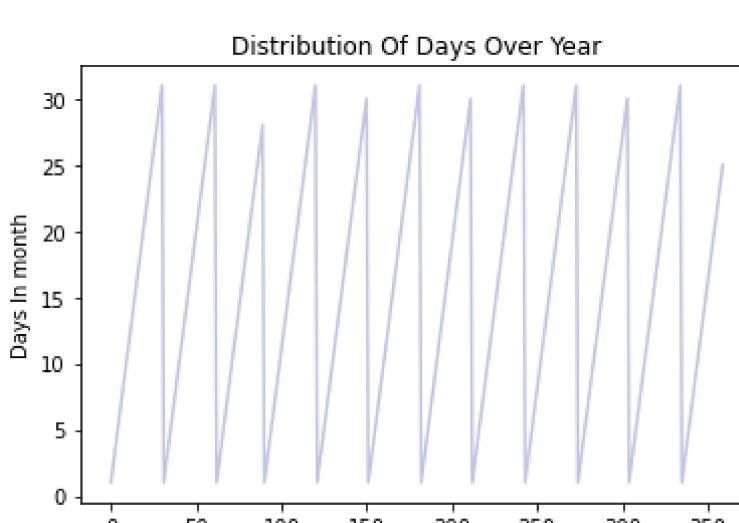
data.head()
```

```
Out[8]:   Date Location MinTemp MaxTemp Rainfall Evaporation Sunshine WindGustDir WindGustSpeed WindDir9am ... Temp3pm RainTo
0 2008-12-01 Albury     13.4     22.9      0.6        NaN        NaN          W         44.0          W ...       21.8
1 2008-12-02 Albury      7.4     25.1      0.0        NaN        NaN        WNW         44.0        NNW ...       24.3
2 2008-12-03 Albury     12.9     25.7      0.0        NaN        NaN        WSW         46.0          W ...       23.2
3 2008-12-04 Albury      9.2     28.0      0.0        NaN        NaN          NE         24.0          SE ...       26.5
4 2008-12-05 Albury     17.5     32.3      1.0        NaN        NaN          W         41.0        ENE ...       29.7
```

5 rows × 30 columns

```
In [9]: # A time period of approximately one year.
section = data[:360]
tm = section["day"].plot(color="#C2C4E2")
tm.set_title("Distribution Of Days Over Year")
tm.set_ylabel("Days In month")
tm.set_xlabel("Days In Year")
```

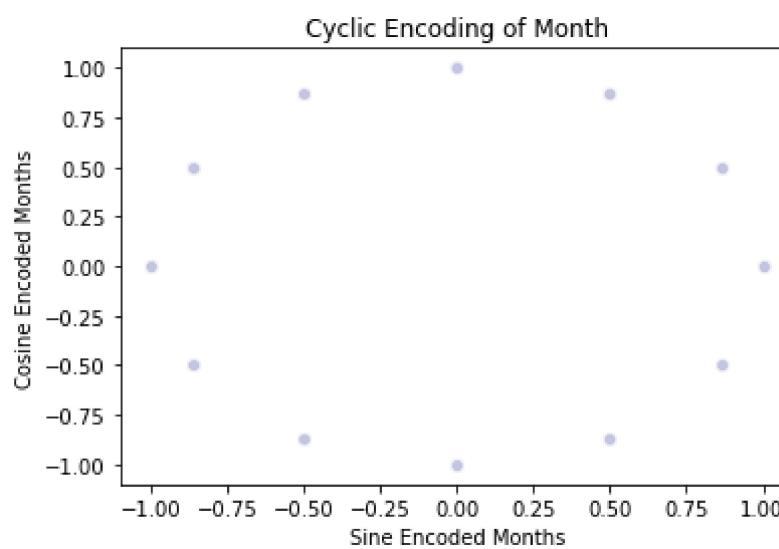
```
Out[9]: Text(0.5, 0, 'Days In Year')
```



As expected, the 'year' attribute in the data repeats, but it doesn't exhibit a continuous cyclical nature. By breaking down months and days into sine and cosine combinations, a continuous cyclic feature is established. These features can be utilized as inputs for an Artificial Neural Network (ANN).

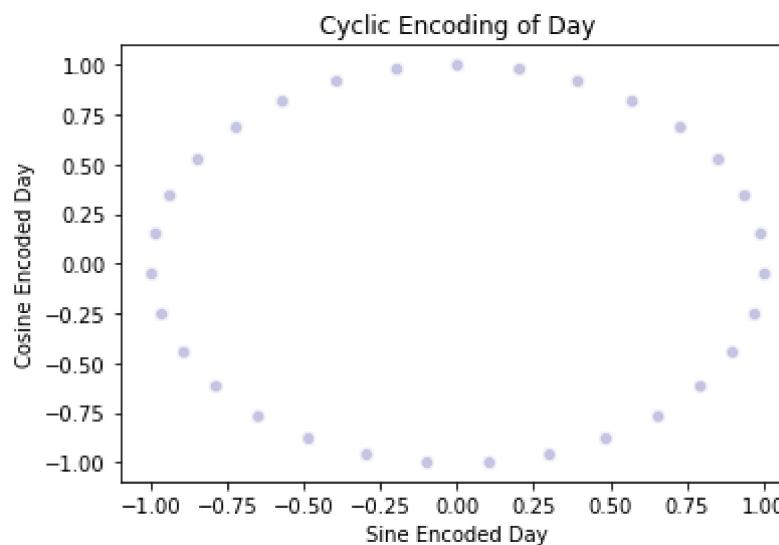
```
In [10]: cyclic_month = sns.scatterplot(x="month_sin",y="month_cos",data=data, color="#C2C4E2")
cyclic_month.set_title("Cyclic Encoding of Month")
cyclic_month.set_ylabel("Cosine Encoded Months")
cyclic_month.set_xlabel("Sine Encoded Months")
```

```
Out[10]: Text(0.5, 0, 'Sine Encoded Months')
```



```
In [11]: cyclic_day = sns.scatterplot(x='day_sin',y='day_cos',data=data, color="#C2C4E2")
cyclic_day.set_title("Cyclic Encoding of Day")
cyclic_day.set_ylabel("Cosine Encoded Day")
cyclic_day.set_xlabel("Sine Encoded Day")
```

```
Out[11]: Text(0.5, 0, 'Sine Encoded Day')
```



Subsequently, I'll address missing values in categorical and numeric attributes separately.

## Categorical variables

Imputing missing values using the mode of the column.

```
In [12]: # Get List of categorical variables
s = (data.dtypes == "object")
object_cols = list(s[s].index)

print("Categorical variables:")
print(object_cols)
```

```
Categorical variables:
['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']
```

```
In [13]: # Missing values in categorical variables
```

```
for i in object_cols:
    print(i, data[i].isnull().sum())
```

```
Location 0
WindGustDir 10326
WindDir9am 10566
WindDir3pm 4228
RainToday 3261
RainTomorrow 3267
```

```
In [ ]: # Imputing missing values with the mode of the column.
```

```
for i in object_cols:
    data[i].fillna(data[i].mode()[0], inplace=True)
```

## Numerical variables

- Filling missing values with median of the column value

```
In [15]: # Get List of numeric variables
t = (data.dtypes == "float64")
num_cols = list(t[t].index)

print("Numeric variables:")
print(num_cols)
```

```
Numeric variables:
['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am', 'Temp3pm', 'month_sin', 'month_cos', 'day_sin', 'day_cos']
```

```
In [16]: # Missing values in numeric variables
```

```

for i in num_cols:
    print(i, data[i].isnull().sum())

MinTemp 1485
MaxTemp 1261
Rainfall 3261
Evaporation 62790
Sunshine 69835
WindGustSpeed 10263
WindSpeed9am 1767
WindSpeed3pm 3062
Humidity9am 2654
Humidity3pm 4507
Pressure9am 15065
Pressure3pm 15028
Cloud9am 55888
Cloud3pm 59358
Temp9am 1767
Temp3pm 3609
month_sin 0
month_cos 0
day_sin 0
day_cos 0

```

In [17]: # Replacing missing values with the median of the column.

```

for i in num_cols:
    data[i].fillna(data[i].median(), inplace=True)

data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
0   Date         145460 non-null   datetime64[ns]
1   Location     145460 non-null   object  
2   MinTemp      145460 non-null   float64 
3   MaxTemp      145460 non-null   float64 
4   Rainfall     145460 non-null   float64 
5   Evaporation  145460 non-null   float64 
6   Sunshine     145460 non-null   float64 
7   WindGustDir  145460 non-null   object  
8   WindGustSpeed 145460 non-null   float64 
9   WindDir9am   145460 non-null   object  
10  WindDir3pm   145460 non-null   object  
11  WindSpeed9am 145460 non-null   float64 
12  WindSpeed3pm 145460 non-null   float64 
13  Humidity9am  145460 non-null   float64 
14  Humidity3pm  145460 non-null   float64 
15  Pressure9am  145460 non-null   float64 
16  Pressure3pm  145460 non-null   float64 
17  Cloud9am     145460 non-null   float64 
18  Cloud3pm     145460 non-null   float64 
19  Temp9am      145460 non-null   float64 
20  Temp3pm      145460 non-null   float64 
21  RainToday    145460 non-null   object  
22  RainTomorrow 145460 non-null   object  
23  year          145460 non-null   int64  
24  month         145460 non-null   int64  
25  month_sin    145460 non-null   float64 
26  month_cos    145460 non-null   float64 
27  day           145460 non-null   int64  
28  day_sin       145460 non-null   float64 
29  day_cos       145460 non-null   float64 

dtypes: datetime64[ns](1), float64(20), int64(3), object(6)
memory usage: 33.3+ MB

```

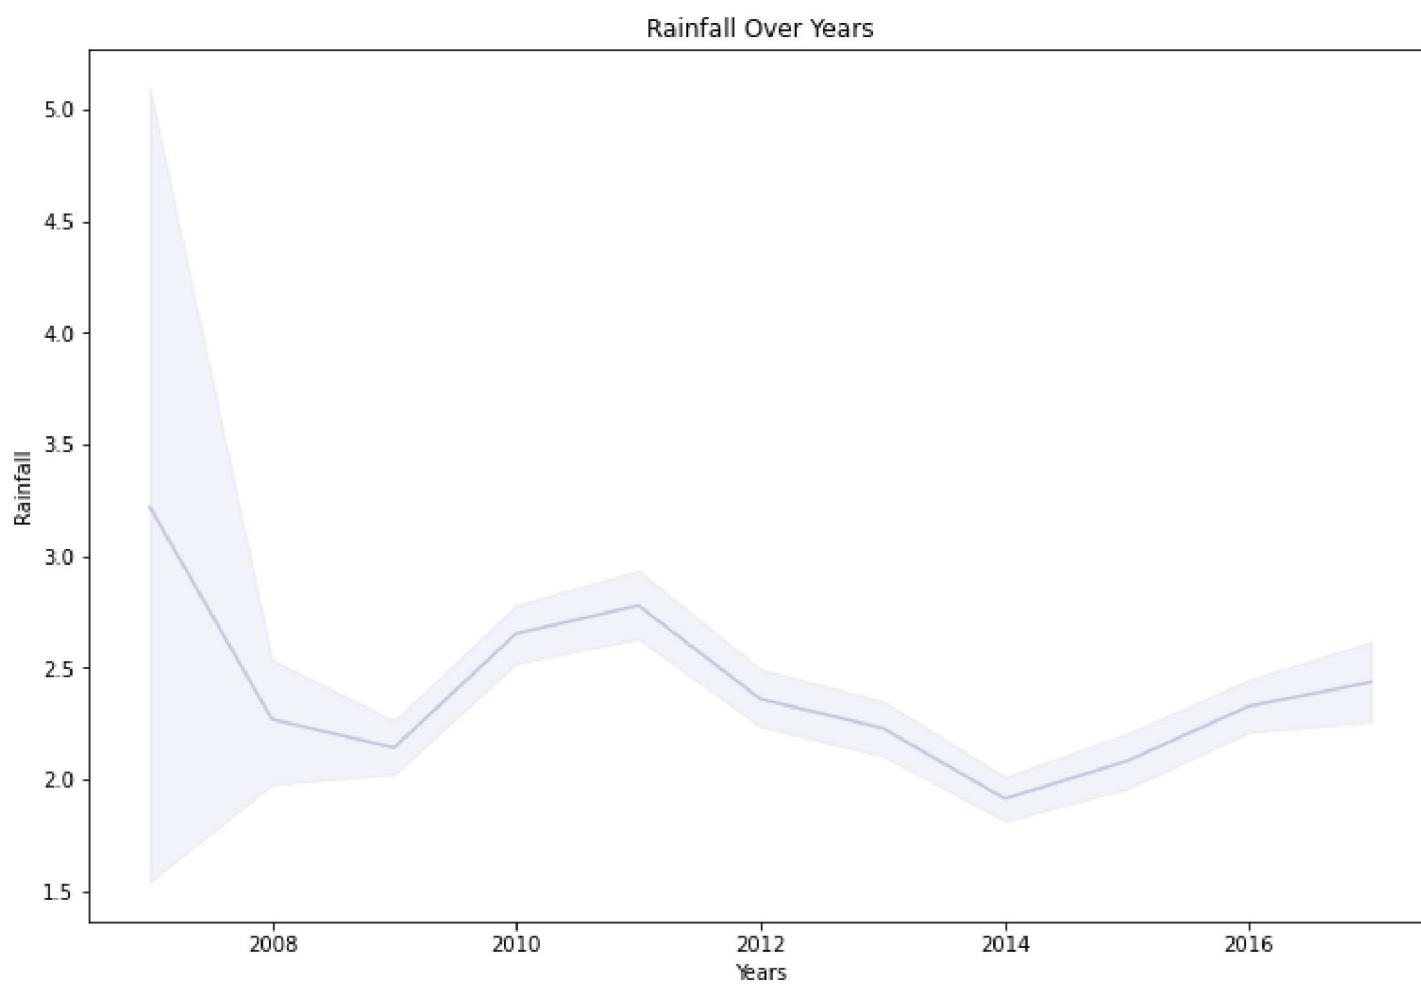
In [18]: #plotting a lineplot rainfall over years

```

plt.figure(figsize=(12,8))
Time_series=sns.lineplot(x=data['Date'].dt.year,y="Rainfall",data=data,color="#C2C4E2")
Time_series.set_title("Rainfall Over Years")
Time_series.set_ylabel("Rainfall")
Time_series.set_xlabel("Years")

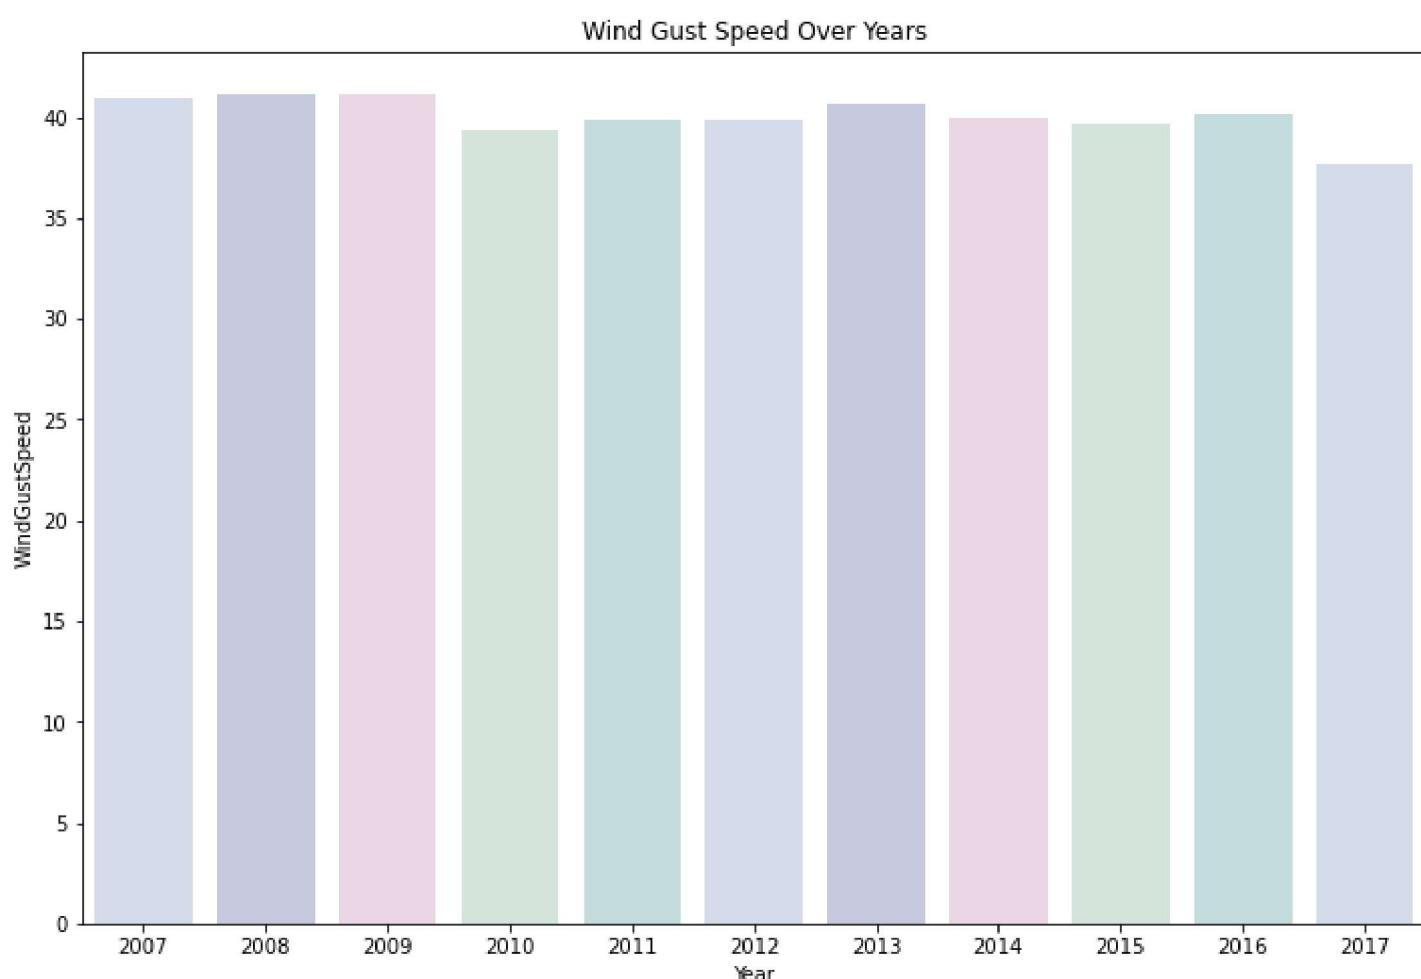
```

Out[18]: Text(0.5, 0, 'Years')



```
In [19]: #Evaluating Wind gust speed over years
colours = ["#D0DBEE", "#C2C4E2", "#EED4E5", "#D1E6DC", "#BDE2E2"]
plt.figure(figsize=(12,8))
Days_of_week=sns.barplot(x=data['Date'].dt.year,y="WindGustSpeed",data=data, ci =None,palette = colours)
Days_of_week.set_title("Wind Gust Speed Over Years")
Days_of_week.set_ylabel("WindGustSpeed")
Days_of_week.set_xlabel("Year")
```

```
Out[19]: Text(0.5, 0, 'Year')
```



#### Data Preprocessing

Data Preprocessing Steps:

Encode columns with categorical data using label encoding

Scale the features

Detect outliers

Remove outliers based on data analysis

#### Label encoding the categorical variable

```
In [20]: # Apply a Label encoder to each column containing categorical data.
label_encoder = LabelEncoder()
for i in object_cols:
    data[i] = label_encoder.fit_transform(data[i])

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Date        145460 non-null   datetime64[ns]
 1   Location    145460 non-null   int64  
 2   MinTemp     145460 non-null   float64
 3   MaxTemp     145460 non-null   float64
 4   Rainfall    145460 non-null   float64
 5   Evaporation 145460 non-null   float64
 6   Sunshine    145460 non-null   float64
 7   WindGustDir 145460 non-null   int64  
 8   WindGustSpeed 145460 non-null   float64
 9   WindDir9am   145460 non-null   int64  
 10  WindDir3pm   145460 non-null   int64  
 11  WindSpeed9am 145460 non-null   float64
 12  WindSpeed3pm 145460 non-null   float64
 13  Humidity9am  145460 non-null   float64
 14  Humidity3pm  145460 non-null   float64
 15  Pressure9am  145460 non-null   float64
 16  Pressure3pm  145460 non-null   float64
 17  Cloud9am    145460 non-null   float64
 18  Cloud3pm    145460 non-null   float64
 19  Temp9am     145460 non-null   float64
 20  Temp3pm     145460 non-null   float64
 21  RainToday   145460 non-null   int64  
 22  RainTomorrow 145460 non-null   int64  
 23  year         145460 non-null   int64  
 24  month        145460 non-null   int64  
 25  month_sin   145460 non-null   float64
 26  month_cos   145460 non-null   float64
 27  day          145460 non-null   int64  
 28  day_sin     145460 non-null   float64
 29  day_cos     145460 non-null   float64
dtypes: datetime64[ns](1), float64(20), int64(9)
memory usage: 33.3 MB
```

```
In [21]: # Preparing attributes of scale data

features = data.drop(['RainTomorrow', 'Date', 'day', 'month'], axis=1) # dropping target and extra columns

target = data['RainTomorrow']

#Set up a standard scaler for the features
col_names = list(features.columns)
s_scaler = preprocessing.StandardScaler()
features = s_scaler.fit_transform(features)
features = pd.DataFrame(features, columns=col_names)

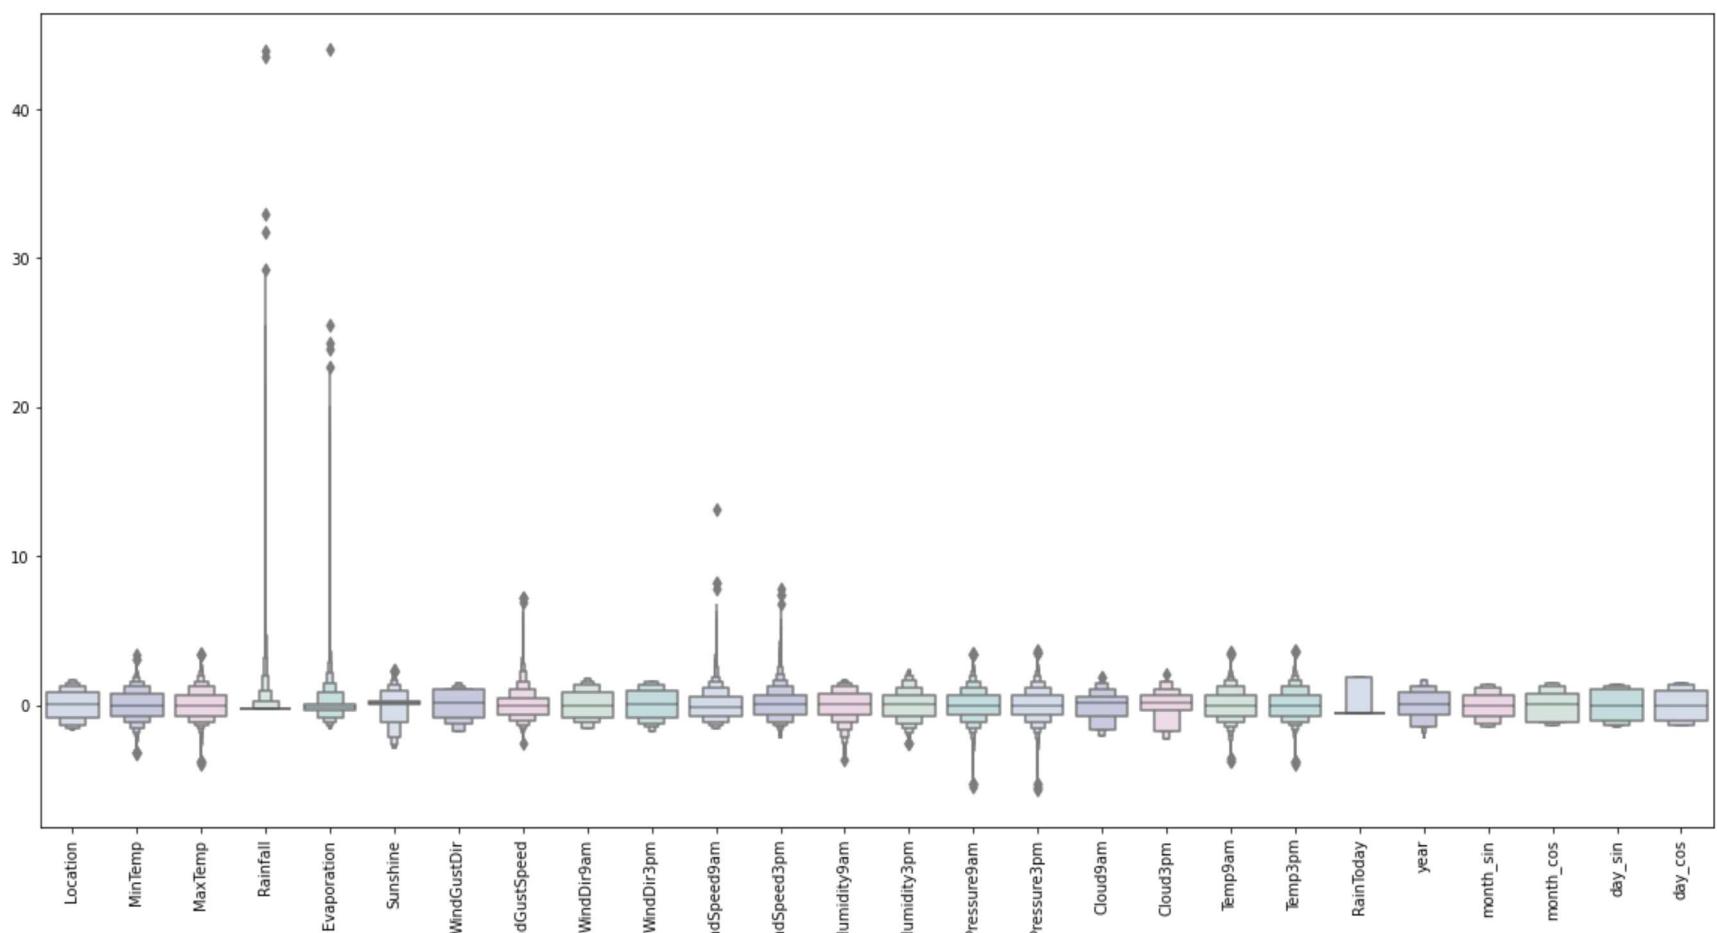
features.describe().T
```

Out[21]:

	count	mean	std	min	25%	50%	75%	max
<b>Location</b>	145460.0	-5.633017e-14	1.000003	-1.672228	-0.899139	0.014511	0.857881	1.701250
<b>MinTemp</b>	145460.0	-4.243854e-15	1.000003	-3.250525	-0.705659	-0.030170	0.723865	3.410112
<b>MaxTemp</b>	145460.0	6.513740e-16	1.000003	-3.952405	-0.735852	-0.086898	0.703133	3.510563
<b>Rainfall</b>	145460.0	9.152711e-15	1.000003	-0.275097	-0.275097	-0.275097	-0.203581	43.945571
<b>Evaporation</b>	145460.0	1.352327e-14	1.000003	-1.629472	-0.371139	-0.119472	0.006361	43.985108
<b>Sunshine</b>	145460.0	-4.338304e-15	1.000003	-2.897217	0.076188	0.148710	0.257494	2.360634
<b>WindGustDir</b>	145460.0	1.864381e-14	1.000003	-1.724209	-0.872075	0.193094	1.045228	1.471296
<b>WindGustSpeed</b>	145460.0	-1.167921e-14	1.000003	-2.588407	-0.683048	-0.073333	0.460168	7.243246
<b>WindDir9am</b>	145460.0	-7.433272e-15	1.000003	-1.550000	-0.885669	0.000105	0.885879	1.771653
<b>WindDir3pm</b>	145460.0	1.791486e-15	1.000003	-1.718521	-0.837098	0.044324	0.925747	1.586813
<b>WindSpeed9am</b>	145460.0	-3.422029e-14	1.000003	-1.583291	-0.793380	-0.116314	0.560752	13.086472
<b>WindSpeed3pm</b>	145460.0	1.618238e-14	1.000003	-2.141841	-0.650449	0.037886	0.611499	7.839016
<b>Humidity9am</b>	145460.0	-4.803490e-15	1.000003	-3.654212	-0.631189	0.058273	0.747734	1.649338
<b>Humidity3pm</b>	145460.0	-6.041889e-15	1.000003	-2.518329	-0.710918	0.021816	0.656852	2.366565
<b>Pressure9am</b>	145460.0	2.313398e-14	1.000003	-5.520544	-0.616005	-0.006653	0.617561	3.471111
<b>Pressure3pm</b>	145460.0	4.709575e-15	1.000003	-5.724832	-0.622769	-0.007520	0.622735	3.653960
<b>Cloud9am</b>	145460.0	-2.525820e-14	1.000003	-2.042425	-0.727490	0.149133	0.587445	1.902380
<b>Cloud3pm</b>	145460.0	4.796901e-15	1.000003	-2.235619	-0.336969	0.137693	0.612356	2.036343
<b>Temp9am</b>	145460.0	-3.332880e-15	1.000003	-3.750358	-0.726764	-0.044517	0.699753	3.599302
<b>Temp3pm</b>	145460.0	-2.901899e-15	1.000003	-3.951301	-0.725322	-0.083046	0.661411	3.653834
<b>RainToday</b>	145460.0	1.263303e-14	1.000003	-0.529795	-0.529795	-0.529795	-0.529795	1.887521
<b>year</b>	145460.0	1.663818e-14	1.000003	-2.273637	-0.697391	0.090732	0.878855	1.666978
<b>month_sin</b>	145460.0	1.653870e-15	1.000003	-1.434333	-0.725379	-0.016425	0.692529	1.401483
<b>month_cos</b>	145460.0	4.043483e-16	1.000003	-1.388032	-1.198979	0.023080	0.728636	1.434192
<b>day_sin</b>	145460.0	-1.982159e-17	1.000003	-1.403140	-1.019170	-0.003198	1.012774	1.396744
<b>day_cos</b>	145460.0	-9.540621e-19	1.000003	-1.392587	-1.055520	-0.044639	1.011221	1.455246

In [22]:

```
#Detecting outliers
#Looking at the scaled features
colours = ["#D0DBEE", "#C2C4E2", "#EED4E5", "#D1E6DC", "#BDE2E2"]
plt.figure(figsize=(20,10))
sns.boxenplot(data = features,palette = colours)
plt.xticks(rotation=90)
plt.show()
```



In [23]:

```
#full data for
features["RainTomorrow"] = target

#Dropping with outlier

features = features[(features["MinTemp"]<2.3)&(features["MinTemp"]>-2.3)]
features = features[(features["MaxTemp"]<2.3)&(features["MaxTemp"]>-2)]
features = features[(features["Rainfall"]<4.5)]
```

```

features = features[(features["Evaporation"]<2.8)]
features = features[(features["Sunshine"]<2.1)]
features = features[(features["WindGustSpeed"]<4)&(features["WindGustSpeed"]>-4)]
features = features[(features["WindSpeed9am"]<4)]
features = features[(features["WindSpeed3pm"]<2.5)]
features = features[(features["Humidity9am"]>-3)]
features = features[(features["Humidity3pm"]>-2.2)]
features = features[(features["Pressure9am"]< 2)&(features["Pressure9am"]>-2.7)]
features = features[(features["Pressure3pm"]< 2)&(features["Pressure3pm"]>-2.7)]
features = features[(features["Cloud9am"]<1.8)]
features = features[(features["Cloud3pm"]<2)]
features = features[(features["Temp9am"]<2.3)&(features["Temp9am"]>-2)]
features = features[(features["Temp3pm"]<2.3)&(features["Temp3pm"]>-2)]

```

features.shape

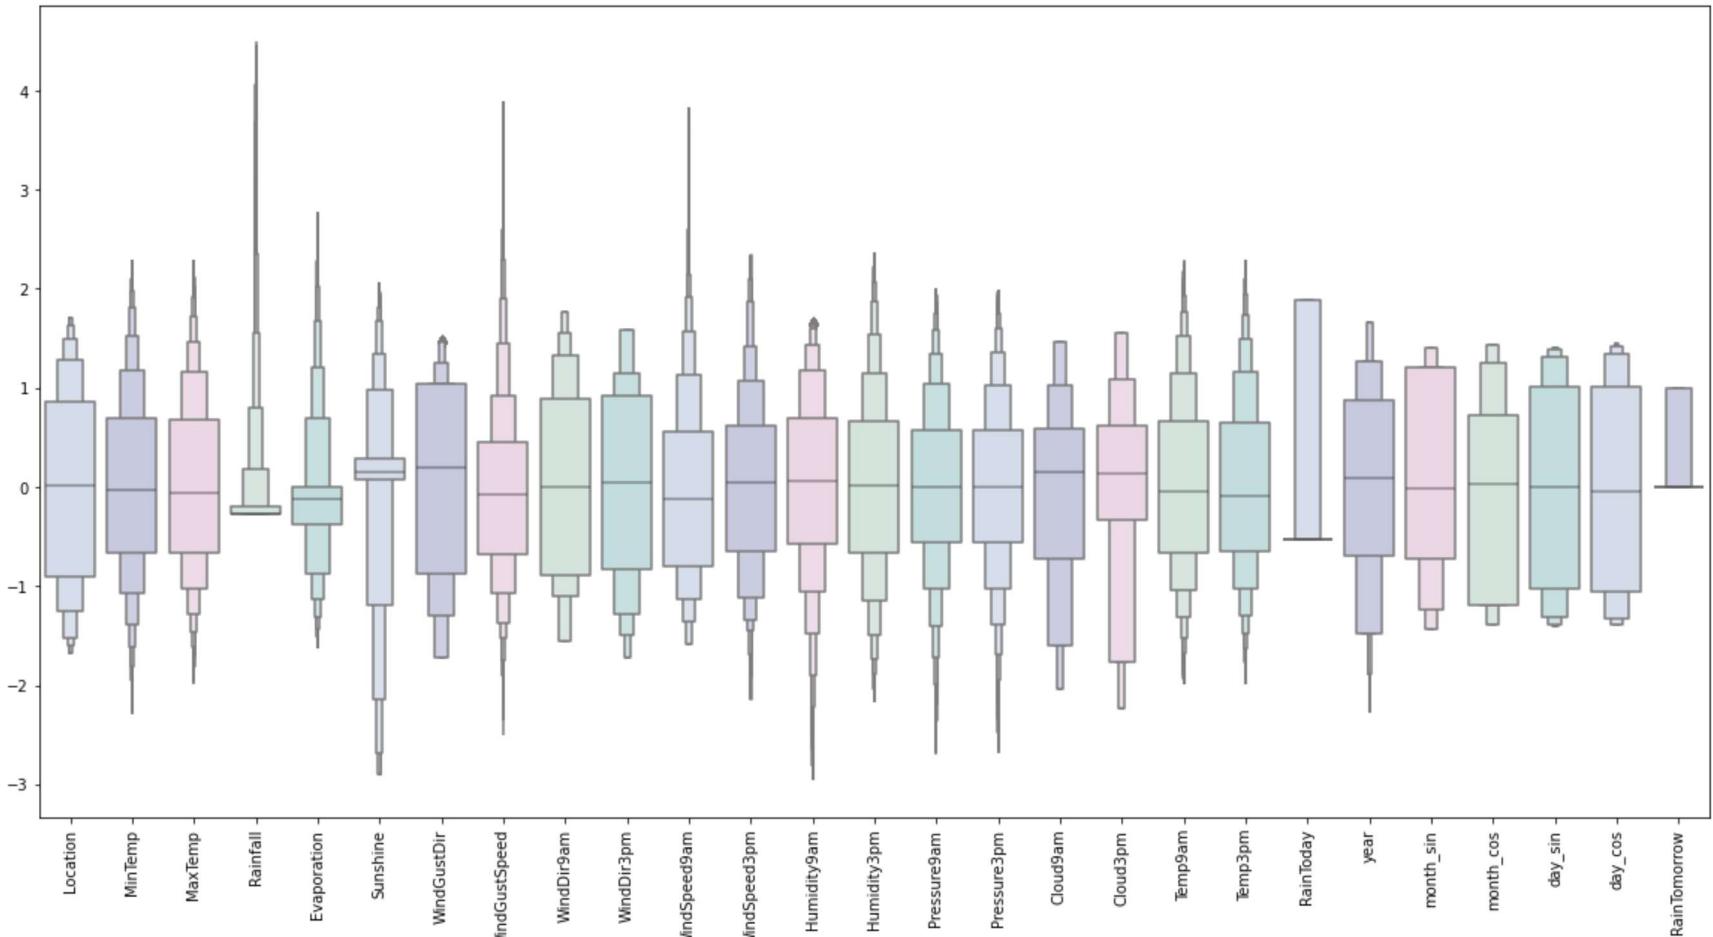
Out[23]: (127536, 27)

In [24]: #Looking at the scaled features without outliers

```

plt.figure(figsize=(20,10))
sns.boxenplot(data = features,palette = colours)
plt.xticks(rotation=90)
plt.show()

```



## Model Building

Steps Involved in Model Building:

Assigning X and y to represent attributes and target labels, respectively.

Splitting the dataset into training and testing sets.

Initializing the neural network.

Defining the architecture by adding layers to the network.

Compiling the neural network.

Training the neural network.

```

In [25]: X = features.drop(["RainTomorrow"], axis=1)
y = features["RainTomorrow"]

# Splitting test and training sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)

X.shape

```

Out[25]: (127536, 26)

```

In [26]: #Early stopping
early_stopping = callbacks.EarlyStopping(
    min_delta=0.001, # minimum amount of change to count as an improvement
    patience=20, # how many epochs to wait before stopping
    restore_best_weights=True,
)

# Initialising the ANN
model = Sequential()

```

```
# Layers

model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu', input_dim = 26))
model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.25))
model.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))

# Compiling the ANN
opt = Adam(learning_rate=0.00009)
model.compile(optimizer = opt, loss = 'binary_crossentropy', metrics = ['accuracy'])

# Train the ANN
history = model.fit(X_train, y_train, batch_size = 32, epochs = 150, callbacks=[early_stopping], validation_split=0.2)
```

Epoch 1/150  
2551/2551 [=====] - 8s 3ms/step - loss: 0.5613 - accuracy: 0.7848 - val\_loss: 0.3887 - val\_accuracy: 0.8245  
Epoch 2/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.4118 - accuracy: 0.8119 - val\_loss: 0.3769 - val\_accuracy: 0.8381  
Epoch 3/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3979 - accuracy: 0.8283 - val\_loss: 0.3695 - val\_accuracy: 0.8406  
Epoch 4/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3978 - accuracy: 0.8293 - val\_loss: 0.3663 - val\_accuracy: 0.8426  
Epoch 5/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3945 - accuracy: 0.8290 - val\_loss: 0.3643 - val\_accuracy: 0.8432  
Epoch 6/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3938 - accuracy: 0.8285 - val\_loss: 0.3637 - val\_accuracy: 0.8447  
Epoch 7/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3897 - accuracy: 0.8333 - val\_loss: 0.3621 - val\_accuracy: 0.8447  
Epoch 8/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3891 - accuracy: 0.8283 - val\_loss: 0.3614 - val\_accuracy: 0.8448  
Epoch 9/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3853 - accuracy: 0.8334 - val\_loss: 0.3605 - val\_accuracy: 0.8452  
Epoch 10/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3879 - accuracy: 0.8310 - val\_loss: 0.3601 - val\_accuracy: 0.8449  
Epoch 11/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3838 - accuracy: 0.8331 - val\_loss: 0.3594 - val\_accuracy: 0.8446  
Epoch 12/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3824 - accuracy: 0.8346 - val\_loss: 0.3580 - val\_accuracy: 0.8454  
Epoch 13/150  
2551/2551 [=====] - 7s 3ms/step - loss: 0.3881 - accuracy: 0.8298 - val\_loss: 0.3580 - val\_accuracy: 0.8444  
Epoch 14/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3878 - accuracy: 0.8297 - val\_loss: 0.3572 - val\_accuracy: 0.8447  
Epoch 15/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3823 - accuracy: 0.8324 - val\_loss: 0.3568 - val\_accuracy: 0.8448  
Epoch 16/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3807 - accuracy: 0.8332 - val\_loss: 0.3568 - val\_accuracy: 0.8433  
Epoch 17/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3846 - accuracy: 0.8310 - val\_loss: 0.3561 - val\_accuracy: 0.8443  
Epoch 18/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3798 - accuracy: 0.8331 - val\_loss: 0.3559 - val\_accuracy: 0.8433  
Epoch 19/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3853 - accuracy: 0.8298 - val\_loss: 0.3550 - val\_accuracy: 0.8447  
Epoch 20/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3802 - accuracy: 0.8316 - val\_loss: 0.3552 - val\_accuracy: 0.8442  
Epoch 21/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3798 - accuracy: 0.8325 - val\_loss: 0.3548 - val\_accuracy: 0.8438  
Epoch 22/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3817 - accuracy: 0.8311 - val\_loss: 0.3549 - val\_accuracy: 0.8434  
Epoch 23/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3790 - accuracy: 0.8330 - val\_loss: 0.3548 - val\_accuracy: 0.8439  
Epoch 24/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3809 - accuracy: 0.8328 - val\_loss: 0.3551 - val\_accuracy: 0.8438  
Epoch 25/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3777 - accuracy: 0.8304 - val\_loss: 0.3555 - val\_accuracy: 0.8433  
Epoch 26/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3784 - accuracy: 0.8310 - val\_loss: 0.3552 - val\_accuracy: 0.8435  
Epoch 27/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3751 - accuracy: 0.8337 - val\_loss: 0.3540 - val\_accuracy: 0.8445  
Epoch 28/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3744 - accuracy: 0.8340 - val\_loss: 0.3539 - val\_accuracy: 0.8437  
Epoch 29/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3769 - accuracy: 0.8340 - val\_loss: 0.3543 - val\_accuracy: 0.8434  
Epoch 30/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3764 - accuracy: 0.8325 - val\_loss: 0.3534 - val\_accuracy: 0.8436  
Epoch 31/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3783 - accuracy: 0.8313 - val\_loss: 0.3537 - val\_accuracy: 0.8439  
Epoch 32/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3760 - accuracy: 0.8333 - val\_loss: 0.3535 - val\_accuracy:

uracy: 0.8437  
Epoch 33/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3732 - accuracy: 0.8349 - val\_loss: 0.3533 - val\_acc  
uracy: 0.8433  
Epoch 34/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3769 - accuracy: 0.8318 - val\_loss: 0.3537 - val\_acc  
uracy: 0.8430  
Epoch 35/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3754 - accuracy: 0.8324 - val\_loss: 0.3521 - val\_acc  
uracy: 0.8448  
Epoch 36/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3787 - accuracy: 0.8298 - val\_loss: 0.3525 - val\_acc  
uracy: 0.8446  
Epoch 37/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3747 - accuracy: 0.8350 - val\_loss: 0.3530 - val\_acc  
uracy: 0.8449  
Epoch 38/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3787 - accuracy: 0.8309 - val\_loss: 0.3519 - val\_acc  
uracy: 0.8448  
Epoch 39/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3757 - accuracy: 0.8313 - val\_loss: 0.3534 - val\_acc  
uracy: 0.8432  
Epoch 40/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3780 - accuracy: 0.8325 - val\_loss: 0.3522 - val\_acc  
uracy: 0.8448  
Epoch 41/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3736 - accuracy: 0.8337 - val\_loss: 0.3515 - val\_acc  
uracy: 0.8467  
Epoch 42/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3749 - accuracy: 0.8336 - val\_loss: 0.3518 - val\_acc  
uracy: 0.8456  
Epoch 43/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3696 - accuracy: 0.8331 - val\_loss: 0.3523 - val\_acc  
uracy: 0.8455  
Epoch 44/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3753 - accuracy: 0.8312 - val\_loss: 0.3524 - val\_acc  
uracy: 0.8450  
Epoch 45/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3723 - accuracy: 0.8344 - val\_loss: 0.3520 - val\_acc  
uracy: 0.8453  
Epoch 46/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3784 - accuracy: 0.8312 - val\_loss: 0.3523 - val\_acc  
uracy: 0.8448  
Epoch 47/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3693 - accuracy: 0.8351 - val\_loss: 0.3534 - val\_acc  
uracy: 0.8440  
Epoch 48/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3720 - accuracy: 0.8343 - val\_loss: 0.3510 - val\_acc  
uracy: 0.8458  
Epoch 49/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3753 - accuracy: 0.8321 - val\_loss: 0.3517 - val\_acc  
uracy: 0.8453  
Epoch 50/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3769 - accuracy: 0.8313 - val\_loss: 0.3521 - val\_acc  
uracy: 0.8456  
Epoch 51/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3715 - accuracy: 0.8356 - val\_loss: 0.3520 - val\_acc  
uracy: 0.8455  
Epoch 52/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3721 - accuracy: 0.8345 - val\_loss: 0.3519 - val\_acc  
uracy: 0.8455  
Epoch 53/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3723 - accuracy: 0.8347 - val\_loss: 0.3528 - val\_acc  
uracy: 0.8450  
Epoch 54/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3764 - accuracy: 0.8311 - val\_loss: 0.3510 - val\_acc  
uracy: 0.8465  
Epoch 55/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3711 - accuracy: 0.8340 - val\_loss: 0.3512 - val\_acc  
uracy: 0.8469  
Epoch 56/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3723 - accuracy: 0.8328 - val\_loss: 0.3510 - val\_acc  
uracy: 0.8463  
Epoch 57/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3759 - accuracy: 0.8309 - val\_loss: 0.3511 - val\_acc  
uracy: 0.8459  
Epoch 58/150  
2551/2551 [=====] - 7s 3ms/step - loss: 0.3715 - accuracy: 0.8345 - val\_loss: 0.3514 - val\_acc  
uracy: 0.8457  
Epoch 59/150  
2551/2551 [=====] - 6s 3ms/step - loss: 0.3708 - accuracy: 0.8351 - val\_loss: 0.3503 - val\_acc  
uracy: 0.8470  
Epoch 60/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3715 - accuracy: 0.8338 - val\_loss: 0.3504 - val\_acc  
uracy: 0.8471  
Epoch 61/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3736 - accuracy: 0.8339 - val\_loss: 0.3506 - val\_acc  
uracy: 0.8469  
Epoch 62/150  
2551/2551 [=====] - 6s 2ms/step - loss: 0.3676 - accuracy: 0.8352 - val\_loss: 0.3506 - val\_acc  
uracy: 0.8461  
Epoch 63/150  
2551/2551 [=====] - 7s 3ms/step - loss: 0.3686 - accuracy: 0.8351 - val\_loss: 0.3508 - val\_acc  
uracy: 0.8462  
Epoch 64/150

```

2551/2551 [=====] - 6s 3ms/step - loss: 0.3712 - accuracy: 0.8341 - val_loss: 0.3518 - val_accuracy: 0.8447
Epoch 65/150
2551/2551 [=====] - 6s 2ms/step - loss: 0.3720 - accuracy: 0.8318 - val_loss: 0.3507 - val_accuracy: 0.8460
Epoch 66/150
2551/2551 [=====] - 6s 2ms/step - loss: 0.3639 - accuracy: 0.8378 - val_loss: 0.3510 - val_accuracy: 0.8463
Epoch 67/150
2551/2551 [=====] - 6s 2ms/step - loss: 0.3721 - accuracy: 0.8325 - val_loss: 0.3507 - val_accuracy: 0.8459
Epoch 68/150
2551/2551 [=====] - 6s 3ms/step - loss: 0.3659 - accuracy: 0.8375 - val_loss: 0.3506 - val_accuracy: 0.8469

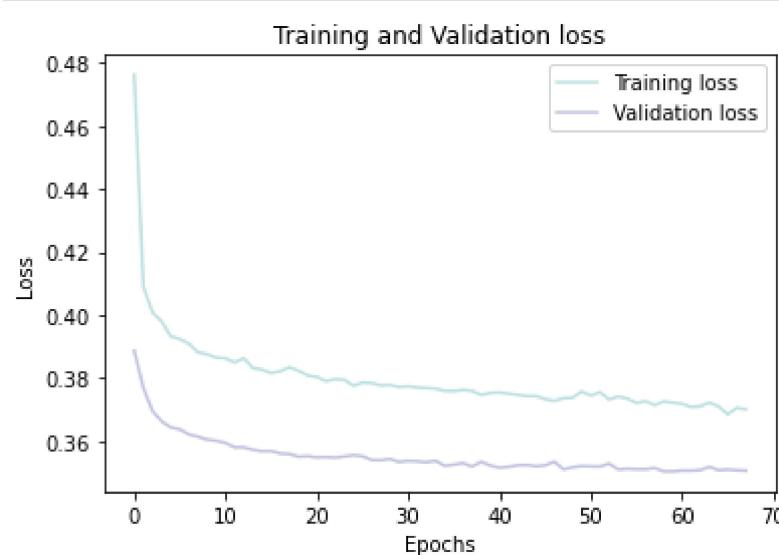
```

Creating a plot to display the training and validation loss across epochs.

```
In [27]: history_df = pd.DataFrame(history.history)

plt.plot(history_df.loc[:, ['loss']], "#BDE2E2", label='Training loss')
plt.plot(history_df.loc[:, ['val_loss']], "#C2C4E2", label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc="best")

plt.show()
```

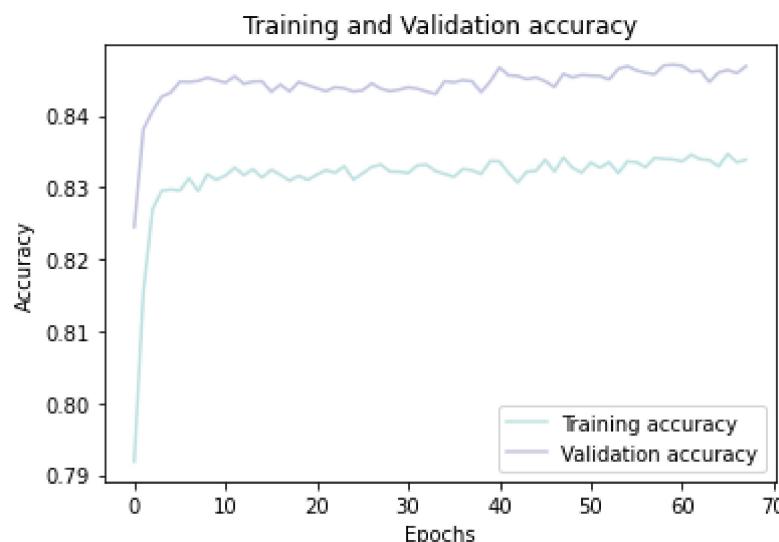


Generating a plot to visualize the training and validation accuracy throughout the epochs.

```
In [28]: history_df = pd.DataFrame(history.history)

plt.plot(history_df.loc[:, ['accuracy']], "#BDE2E2", label='Training accuracy')
plt.plot(history_df.loc[:, ['val_accuracy']], "#C2C4E2", label='Validation accuracy')

plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



## Conclusion

Testing on the test set

Evaluating the confusion matrix

Assessing the classification report

```
In [29]: # Predicting the test set results
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)
```

```
In [30]: # confusion matrix
cmap1 = sns.diverging_palette(260,-10,s=50, l=75, n=5, as_cmap=True)
plt.subplots(figsize=(12,8))
```

```
cf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(cf_matrix/np.sum(cf_matrix), cmap = cmap1, annot = True, annot_kws = {'size':15})
```

Out[30]: <AxesSubplot:>



In [31]: `print(classification_report(y_test, y_pred))`

	precision	recall	f1-score	support
0	0.86	0.96	0.91	20110
1	0.76	0.43	0.55	5398
accuracy			0.85	25508
macro avg	0.81	0.70	0.73	25508
weighted avg	0.84	0.85	0.83	25508